

Jeremy Springman, PhD  
Research Assistant Professor  
University of Pennsylvania  
[jrspringman.github.io](https://jrspringman.github.io)

I am an applied researcher (and longtime Redditor) passionate about bringing advanced experimentation methodologies to real-world decision-making. As a Research Assistant Professor at the University of Pennsylvania, I have spent the last decade using causal inference and machine learning to inform high-stakes decision-making by policymakers in the U.S. government. From designing and analyzing complex experimental studies to delivering rigorous, data-driven insights that inform strategic decisions, I have cultivated the skills and experience required to bring data science excellence to Reddit's two-sided ads marketplace.

Over the past six years, I have managed and mentored an interdisciplinary team of twelve researchers and data scientists. Under my leadership, this team has planned and executed many large-scale projects, ensuring that the findings are both rigorous and actionable. Since 2022, I have been awarded nearly \$3.5 million as a principal investigator on both federal and private research grants, carrying projects from conceptualization and proposal through design, implementation, and the dissemination of results to the scientific and policy communities.

I have deep experience leading randomized experiments and causal inference projects commissioned by stakeholders to generate actionable insights. These projects have given me years of experience using advanced methods to address novel design and measurement challenges. I have designed and publicly pre-registered more than a dozen randomized experiments using bespoke simulation-based power calculations. I have also leveraged tools like non-bipartite matching, matched-group randomization, and re-randomization to extract maximum statistical precision from small and high variance samples and deployed sophisticated approaches to adjust for multiple-hypothesis testing, including Bayesian regularization. I have also designed trials to deal with the reality of social networks by collecting and using network data to estimate the size of treatment spillovers through network connections.

In addition to generating academic publications and contributing to our theoretical understanding of the social world, my research has directly informed policy. As a research associate at Duke University, I led an evaluation that provided rigorous evidence that a U.S. government foreign aid program was failing to accomplish its goals, resulting in the program's cancellation and a more effective use of government resources.

In another project, I worked with policymakers at the U.S. Agency for International Development to design a conjoint survey experiment that answered pressing questions about the strategic behavior of non-profits receiving U.S. support in highly repressive countries. Based on the findings, I produced a policy brief recommending changes to how the agency supports local partners in repressive countries and presented the recommendations to high-level officials. A third project used insights from behavioral psychology and economics to design and evaluate an intervention to increase civic engagement and social cohesion among youth in highly polarized countries. This resulted in a highly effective intervention that my partners are seeking to scale across developing countries.

I also have extensive background in machine learning research. As a post-doctoral research associate, I co-founded a high-profile project to build a public-facing early warning system forecasting political instability in emerging markets. Managing a team of data scientists and PhD students, I led the construction of a research infrastructure to continually update a large, highly multilingual

corpus of high-quality media and use deep learning translation models and fine-tuned transformers and LLMs to extract information from text. This project's funding was renewed over six years and attracted support from the U.S. Department of Defense, Agency for International Development, and private donors.

My lab used data from this project to publish more than a dozen commissioned policy reports, develop public-facing data dashboards averaging 150 hours of monthly active usage time by decision-makers, and attracted 1,200 unique users from 70 countries over a 10 month period, with more than 800 policymakers signing up to receive our monthly reports. The success of this project led to my promotion to Research Assistant Professor, where I manage three full-time data scientists, five pre-doctoral fellows, multiple affiliated faculty and PhD students, and a large pool of undergraduates.

My experience working closely with policymakers and stakeholders prepares me to collaborate with Reddit's product managers and engineers to identify challenges, design solutions, and measure their impact on key outcomes. I am passionate about translating scientific research into action, and have given invited talks to hundreds of policymakers across more than a dozen agencies in five countries. I specialize in communicating technical material to non-technical audiences: from policy officials at the U.S. Department of State to civil society stakeholders across emerging markets, I have built a track record of translating complex analyses into clear, strategic recommendations. Furthermore, my international experience makes me well prepared to work with Reddit's global community of users and advertisers.

I have deep technical proficiency in statistical programming (R and Python), document preparation and visualization (markdown, Quarto, and Shiny), version control (git, GitHub, and renv), programming online surveys (Javascript and html), and have managed teams using high-performance computing environments, NoSQL document-stage databases, and core machine learning Python libraries (scikit-learn, XGBoost, LightGBM, PyTorch, Optuna). I'm excited to deploy these skills in the technology sector, where analytics can inform decisions at an even faster pace.

Working at a university has also given the opportunity to teach and mentor students at the undergraduate and graduate level. My academic career has taught me the importance of mentorship, and I plan to continue fostering data science best practices by guiding junior colleagues and building a collaborative culture that values creativity and rigor.

As an applied researcher, I have been extremely successful at winning highly competitive research grants. However, recent terminations of major federal research initiatives have cut-off opportunities for high-impact collaboration between external researchers and government agencies. I am now looking to transfer my skills to the private sector.

Reddit is an ideal fit for several reasons. Having followed Reddit as an active user for years, I am enthusiastic about contributing to the platform's growth and evolution. As the social media ecosystem becomes more fractured, Reddit's unique role as a community-driven platform makes it a critical online space. The potential to make a tangible impact on millions of users worldwide is an opportunity that I would be honored to embrace. Finally, Reddit's New York office is easily accessible by regional rail from my home in Philadelphia.

In summary, I am highly qualified and extremely excited to accelerate advertiser performance through cutting-edge experimentation, refine key metrics to empower product teams, and advance novel techniques — including ML approaches such as causal forests — within Reddit's advertising technology. Thank you for considering my application.